

Module 1: Pig Latin

Write a Pig script to load the `transaction_data.csv` dataset, filter rows where amount is greater than 300, and store the result in a new file.

```
hadoop fs -put Desktop/transaction_data.csv
```

```
$ pig
```

```
transactions = LOAD 'transaction_data.csv' USING PigStorage(',') AS (transaction_id:chararray,  
user_id:chararray, amount:float, status:chararray);
```

```
filtered_transactions = FILTER transactions BY amount > 300;
```

```
STORE filtered_transactions INTO 'filtered_transactions'
```

```
USING PigStorage(',');
```

```
$ hadoop fs -get filtered_transactions
```

```
$ cat filtered_transactions/part-m-00000
```

OUTPUT

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ hadoop fs -get filtered_transactions;  
[cloudera@quickstart ~]$ cat filtered_transactions/part-m-00000;  
102,2,400.0,Failed  
104,4,600.0,Completed  
105,5,320.0,Failed  
107,2,500.0,Completed  
109,4,620.0,Completed  
110,5,350.0,Completed  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ hadoop fs -put Desktop/transaction_data.csv  
put: 'transaction_data.csv': File exists  
[cloudera@quickstart ~]$ pig  
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).  
log4j:WARN Please initialize the log4j system properly.  
log4j:WARN See http://logging.apache.org/log4j/1.2faq.html#noconfig for more info.  
2024-10-01 23:42:14.864 [main] INFO org.apache.pig.Main - Apache Pig version 0.12.0-cdh5.12.0 (reexported) compile  
d Jun 29 2017. 04:34:31  
2024-10-01 23:42:14.865 [main] INFO org.apache.pig.Main - Logging error messages to: /home/cloudera/pig.172785133  
4795.log  
2024-10-01 23:42:14.919 [main] INFO org.apache.pig.impl.util.Utils - Default bootstrap file /home/cloudera/.pigbootu  
p not found  
2024-10-01 23:42:15.957 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is depr  
ecated. Instead, use mapreduce.jobtracker.address  
2024-10-01 23:42:15.957 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.957 [main] INFO org.apache.pig.backend.hadoop.executionengine.MExecutionEngine - Connecting t  
o hadoop file system at: hdfs://quickstart.cloudera:8020  
2024-10-01 23:42:15.961 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is depr  
ecated. Instead, use mapreduce.jobtracker.address  
2024-10-01 23:42:15.991 [main] INFO org.apache.pig.backend.hadoop.executionengine.MExecutionEngine - Connecting t  
o map-reduce job tracker at: localhost:8021  
2024-10-01 23:42:15.992 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.989 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is depr  
ecated. Instead, use mapreduce.jobtracker.address  
2024-10-01 23:42:15.987 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.954 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.955 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is depr  
ecated. Instead, use mapreduce.jobtracker.address  
2024-10-01 23:42:15.946 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.947 [main] INFO org.apache.pig.backend.hadoop.executionengine.MExecutionEngine - Connecting t  
o map-reduce job tracker at: localhost:8021  
2024-10-01 23:42:15.944 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.945 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is depr  
ecated. Instead, use mapreduce.jobtracker.address  
2024-10-01 23:42:15.931 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.931 [main] INFO org.apache.pig.backend.hadoop.executionengine.MExecutionEngine - Connecting t  
o map-reduce job tracker at: localhost:8021  
2024-10-01 23:42:15.925 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.926 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is depr  
ecated. Instead, use mapreduce.jobtracker.address  
2024-10-01 23:42:15.909 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.909 [main] INFO org.apache.pig.backend.hadoop.executionengine.MExecutionEngine - Connecting t  
o map-reduce job tracker at: localhost:8021  
2024-10-01 23:42:15.908 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is depreca  
ted. Instead, use fs.defaultFS  
2024-10-01 23:42:15.908 [main] INFO org.apache.pig.backend.hadoop.executionengine.MExecutionEngine - Connecting t  
o map-reduce job tracker at: localhost:8021  
[cloudera@quickstart ~]$
```

```

2024-10-01 23:46:44,324 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher
r - More information at: http://localhost:50030/jobdetails.jsp?jobid=job_1727847831054_0001
2024-10-01 23:46:44,473 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher
r - 0% complete
2024-10-01 23:48:11,229 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2024-10-01 23:48:17,629 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduce
s
2024-10-01 23:48:17,982 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2024-10-01 23:48:18,003 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2024-10-01 23:46:26 2024-10-01 23:48:17 FILTER

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReductime A
lias Feature Outputs
job_1727847831054_0001 1 0 23 23 23 23 n/a n/a n/a n/a filtered_transactions,transactions MAP_ONLY h
dfs://quickstart.cloudera:8020/user/cloudera/filtered_transactions,

Input(s):
Successfully read 11 records (613 bytes) from: "hdfs://quickstart.cloudera:8020/user/cloudera/transaction_data.csv"

Output(s):
Successfully stored 6 records (126 bytes) in: "hdfs://quickstart.cloudera:8020/user/cloudera/filtered_transactions"

Counters:
Total records written : 6
Total bytes written : 126
Spillable Memory Manager spill count : 0
Total bugs proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1727847831054_0001

2024-10-01 23:48:18,174 [main] WARN org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Encountered Warning FIELD_DISCARDED_TY
PE CONVERSION_FAILED 1 time(s).
2024-10-01 23:48:18,174 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!

```

2. Use Pig to calculate the average amount from the transaction_data.csv grouped by status. Concatenate user_id and status into a new field in a Pig script.

```
transactions = LOAD 'transaction_data.csv' USING PigStorage(',') AS (transaction_id:chararray,
user_id:chararray, amount:float, status:chararray);
```

```
grouped_by_status = GROUP transactions BY status;
```

```
average_amount = FOREACH grouped_by_status GENERATE group AS status,
AVG(transactions.amount) AS avg_amount;
```

```
STORE average_amount INTO 'average_amount_by_status' USING PigStorage(',');
```

```
$ hadoop fs -get average_amount_by_status
```

```
$ cat average_amount_by_status/part-r-00000
```

```

cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ hadoop fs -get average_amount_by_status;
[cloudera@quickstart ~]$ cat average_amount_by_status/part-r-00000;
Failed,280.0
status,
Completed,394.2857142857143
[cloudera@quickstart ~]$

```

3.Concatenate user_id and status into a new field in a Pig script.

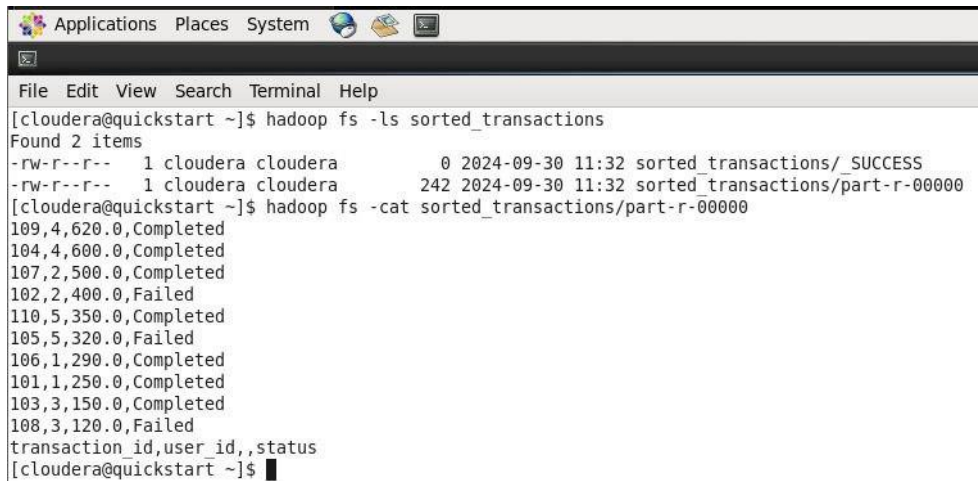
```
transactions_with_concat = FOREACH transactions GENERATE transaction_id, user_id, amount,  
status,
```

```
CONCAT(user_id, status) AS user_status_concat;
```

```
STORE transactions_with_concat INTO 'transactions_with_concat' USING PigStorage(',');
```

```
$ hadoop fs -get transactions_with_concat
```

```
$ cat transactions_with_concat/part-r-00000
```



```
Applications Places System
[cloudera@quickstart ~]$ hadoop fs -ls sorted_transactions
Found 2 items
-rw-r--r-- 1 cloudera cloudera 0 2024-09-30 11:32 sorted_transactions/ SUCCESS
-rw-r--r-- 1 cloudera cloudera 242 2024-09-30 11:32 sorted_transactions/part-r-00000
[cloudera@quickstart ~]$ hadoop fs -cat sorted_transactions/part-r-00000
109,4,620.0,Completed
104,4,600.0,Completed
107,2,500.0,Completed
102,2,400.0,Failed
110,5,350.0,Completed
105,5,320.0,Failed
106,1,290.0,Completed
101,1,250.0,Completed
103,3,150.0,Completed
108,3,120.0,Failed
transaction_id,user_id,,status
[cloudera@quickstart ~]$
```

4.Sort the transaction_data.csv dataset by amount in descending order.

```
sorted_transactions = ORDER transactions BY amount DESC;
```

```
STORE sorted_transactions INTO 'sorted_transactions' USING PigStorage(',');
```

```
$ hadoop fs -get sorted_transactions
```

```
$ cat sorted_transactions/part-r-00000
```

```
Applications Places System
[cloudera@quickstart ~]$ hadoop fs -ls transactions_with_concat
Found 2 items
-rw-r--r-- 1 cloudera cloudera 0 2024-09-30 11:26 transactions_with_concat/_SUCCESS
-rw-r--r-- 1 cloudera cloudera 357 2024-09-30 11:26 transactions_with_concat/part-m-00000
^[[A[cloudera@quickstart ~]$ hadoop fs -cat transactions_with_concat/part-m-00000
transaction_id,user_id,,status,user_idstatus
101,1,250.0,Completed,1Completed
102,2,400.0,Failed,2Failed
103,3,150.0,Completed,3Completed
104,4,600.0,Completed,4Completed
105,5,320.0,Failed,5Failed
106,1,290.0,Completed,1Completed
107,2,500.0,Completed,2Completed
108,3,120.0,Failed,3Failed
109,4,620.0,Completed,4Completed
110,5,350.0,Completed,5Completed
[cloudera@quickstart ~]$
```

5. Use Pig to find the maximum and minimum amount values in the transaction_data.csv.

max_min_amount = FOREACH (GROUP transactions ALL) GENERATE MAX(transactions.amount) AS max_amount, MIN(transactions.amount) AS min_amount;

STORE max_min_amount INTO 'max_min_amount' USING PigStorage(',');

\$ hadoop fs -get max_min_amount

\$ cat max_min_amount/part-r-00000

```
Applications Places System
[cloudera@quickstart ~]$ hadoop fs -ls max_min_amount
Found 2 items
-rw-r--r-- 1 cloudera cloudera 0 2024-10-02 01:14 max_min_amount/_SUCCESS
-rw-r--r-- 1 cloudera cloudera 12 2024-10-02 01:14 max_min_amount/part-r-00000
[cloudera@quickstart ~]$ hadoop fs -cat max_min_amount/part-r-00000
620.0,120.0
[cloudera@quickstart ~]$
```

1. Write a HiveQL query to create a partitioned table on the status column from transaction_data.csv

- hive;
- CREATE DATABASE retail;
- USE retail;
- CREATE TABLE transaction(transaction_id INT,user_id INT, amount FLOAT,status STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE TBLPROPERTIES ("skip.header.line.count"="1");

```
[cloudera@quickstart ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> CREATE DATABASE retail;
OK
Time taken: 4.644 seconds
hive> use retail;
OK
Time taken: 0.24 seconds
hive> CREATE TABLE transaction(
  >   transaction_id INT,
  >   user_id INT,
  >   amount FLOAT,
  >   status STRING
  > )
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ','
  > STORED AS TEXTFILE
  > TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 1.198 seconds
hive> █
```

2. Insert data into a Hive table from the transaction_data.csv dataset.

LOAD DATA LOCAL INPATH 'Desktop/transaction_data.csv' INTO TABLE transaction;

```
hive> LOAD DATA LOCAL INPATH 'Desktop/transaction_data.csv' INTO TABLE transaction;
Loading data to table retail.transaction
Table retail.transaction stats: [numFiles=1, totalSize=228]
OK
Time taken: 4.734 seconds
hive> select * from transaction;
OK
101      1      250.0   Completed
102      2      400.0   Failed
103      3      150.0   Completed
104      4      600.0   Completed
105      5      320.0   Failed
106      1      290.0   Completed
107      2      500.0   Completed
108      3      120.0   Failed
109      4      620.0   Completed
110      5      350.0   Completed
Time taken: 1.626 seconds, Fetched: 10 row(s)
hive> █
```

3. Perform a left outer join between the employee and transaction tables in Hive.

- CREATE TABLE employee(emp_id INT,name STRING,age INT, department STRING, salary INT)
ROW FORMAT DELIMITEDTERMINATED BY ','STORED AS TEXTFILE TBLPROPERTIES
("skip.header.line.count"="1");

```
hive> CREATE TABLE employee(  
  >   emp_id INT,  
  >   name STRING,  
  >   age INT,  
  >   department STRING,  
  >   salary INT  
  > )  
  > ROW FORMAT DELIMITED  
  > FIELDS TERMINATED BY ','  
  > STORED AS TEXTFILE  
  > TBLPROPERTIES ("skip.header.line.count"="1");  
OK  
Time taken: 2.007 seconds
```

- LOAD DATA LOCAL INPATH 'Desktop/employee_data.csv' INTO TABLE employee;

```
hive> LOAD DATA LOCAL INPATH 'Desktop/employee_data.csv' INTO TABLE employee;  
Loading data to table default.employee  
Table default.employee stats: [numFiles=1, totalSize=244]  
OK  
Time taken: 1.56 seconds  
hive> set mapreduce.map.memory.mb=4096;  
hive> set mapreduce.reduce.memory.mb=4096;  
hive> set hive.auto.convert.join=false;
```

- set mapreduce.map.memory.mb=4096;
- set mapreduce.reduce.memory.mb=4096;
- set hive.auto.convert.join=false;
- SELECT * FROM employee e LEFT OUTER JOIN transaction t ON e.emp_id = t.user_id;

```

File Edit View Search Terminal Help
hive> SELECT *
  > FROM employee e
  > LEFT OUTER JOIN transaction t
  > ON e.emp_id = t.user_id;
Query ID = cloudera_20241001232323_f230ab10-8cc0-4544-b6ff-64c857158177
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1727846571552_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1727846571552_0002/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1727846571552_0002
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 1
2024-10-01 23:24:25,811 Stage-1 map = 0%, reduce = 0%
2024-10-01 23:24:47,751 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 5.72 sec
2024-10-01 23:24:58,869 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 7.82 sec
2024-10-01 23:25:06,498 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.61 sec
MapReduce Total cumulative CPU time: 10 seconds 610 msec
Ended Job = job_1727846571552_0002
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2 Reduce: 1 Cumulative CPU: 10.61 sec HDFS Read: 14592 HDFS Write: 330 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 610 msec
OK
1 Alice 25 HR 50000 NULL NULL NULL NULL
2 Bob 35 Finance 70000 NULL NULL NULL NULL
3 Charlie 28 HR 60000 NULL NULL NULL NULL
4 David 45 IT 90000 NULL NULL NULL NULL
5 Eva 32 Finance 80000 NULL NULL NULL NULL
6 Frank 29 IT 75000 NULL NULL NULL NULL
7 Grace 40 HR 65000 NULL NULL NULL NULL
8 Hank 38 IT 85000 NULL NULL NULL NULL
9 Ivy 30 Finance 72000 NULL NULL NULL NULL
10 Jack 50 HR 78000 NULL NULL NULL NULL
Time taken: 72.131 seconds, Fetched: 10 row(s)
hive>

```

4. Drop the employee table in Hive if it exists.

DROP TABLE IF EXISTS employee;

```

hive> DROP TABLE IF EXISTS employee;
OK

```

Time taken: 2.927 seconds

```
hive>
```

cloudera@quickstart:~

5. Group the transaction table by status and calculate the total count of records for each group.

SELECT status, COUNT(*) FROM transaction GROUP BY status;

```

hive> SELECT status, COUNT(*) FROM transaction GROUP BY status;
Query ID = cloudera_20241001233636_7505d102-533a-44a3-b652-ee4d3b3b7deb
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1727846571552_0003, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1727846571
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1727846571552_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-10-01 23:37:04,664 Stage-1 map = 0%, reduce = 0%
2024-10-01 23:37:26,785 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.37 sec
2024-10-01 23:37:39,550 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.74 sec
MapReduce Total cumulative CPU time: 7 seconds 740 msec
Ended Job = job_1727846571552_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.74 sec HDFS Read: 7734 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 740 msec
OK
Time taken: 10.830 seconds
hive>

```

```
hive> LOAD DATA LOCAL INPATH 'Desktop/transaction_data.csv' INTO TABLE transaction;
Loading data to table default.transaction
Table default.transaction stats: [numFiles=1, totalSize=228]
OK
Time taken: 5.209 seconds
hive> SELECT status, COUNT(*) FROM transaction GROUP BY status;
Query ID = cloudera_20241002103333_7da1dd62-41d2-43b2-983f-b9230eda6ee3
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1727889282011_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1727889282011_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1727889282011_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2024-10-02 10:34:11,922 Stage-1 map = 0%, reduce = 0%
2024-10-02 10:34:32,397 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.41 sec
2024-10-02 10:34:44,334 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.48 sec
MapReduce Total cumulative CPU time: 7 seconds 480 msec
Ended Job = job_1727889282011_0001
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.48 sec HDFS Read: 7522 HDFS Write: 21 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 480 msec
OK
Completed      7
Failed        3
Time taken: 68.054 seconds, Fetched: 2 row(s)
hive> █
```


Module 3: Spark

1. Select the user_id and amount columns from the transaction_data.csv in Spark.

Step 1: Create the transaction data frame.

```
val df = Seq((101,1,250, "Completed"), (102,2,400,"Failed"), (103,3,150,"Completed"),
(104,4,600,"Completed"), (105,5,320,"Failed"), (106,1,290,"Completed"),
(107,2,500,"Completed"), (108,3,120,"Failed"), (109,4,620,"Completed"),
(110,5,350,"Completed")).toDF("transaction_id","user_id","amount","status")
```

```
scala> val df = Seq((101,1,250, "Completed"), (102,2,400,"Failed"), (103,3,150,"Completed"), (104,4,600,"Completed"), (105,5,320,"Failed"),
(106,1,290,"Completed"), (107,2,500,"Completed"), (108,3,120,"Failed"), (109,4,620,"Completed"), (110,5,350,"Completed")).toDF("transaction_id","user_id","amount","status")
df: org.apache.spark.sql.DataFrame = [transaction_id: int, user_id: int, amount: int, status: string]
```

```
scala> df.show()
+-----+-----+-----+-----+
|transaction_id|user_id|amount|status|
+-----+-----+-----+-----+
|101|1|250|Completed|
|102|2|400|Failed|
|103|3|150|Completed|
|104|4|600|Completed|
|105|5|320|Failed|
|106|1|290|Completed|
|107|2|500|Completed|
|108|3|120|Failed|
|109|4|620|Completed|
|110|5|350|Completed|
+-----+-----+-----+-----+
```

Step 2: Select only the specified columns and display it.

```
val selectedColumns = df.select(df("user_id"), df("amount"))
selectedColumns.show()
```

```
scala> val selectedColumns = df.select(df("user_id"), df("amount"))
selectedColumns: org.apache.spark.sql.DataFrame = [user_id: int, amount: int]
```

```
scala> selectedColumns.show()
```

```
+-----+-----+
|user_id|amount|
+-----+-----+
|1|250|
|2|400|
|3|150|
|4|600|
|5|320|
|1|290|
|2|500|
|3|120|
|4|620|
|5|350|
+-----+-----+
```

```
scala> █
```

2. Sort the DataFrame by amount in descending order.

Step 1: Sort the data frame that you created in the previous question.

```
val sortedDF = df.orderBy(desc("amount"))
```

Step 2: Display.

```
sortedDF.show()
```

```
scala> val sortedDF = df.orderBy(desc("amount"))
sortedDF: org.apache.spark.sql.DataFrame = [transaction_id: int, user_id: int, amount: int, status: string]

scala> sortedDF.show()
+-----+-----+-----+-----+
|transaction_id|user_id|amount|status|
+-----+-----+-----+-----+
|109|4|620|Completed|
|104|4|600|Completed|
|107|2|500|Completed|
|102|2|400|Failed|
|110|5|350|Completed|
|105|5|320|Failed|
|106|1|290|Completed|
|101|1|250|Completed|
|103|3|150|Completed|
|108|3|120|Failed|
+-----+-----+-----+-----+

scala>
```

3. Perform a left join between the transaction and employee DataFrames based on user_id and emp_id.

Step 1: Create the employee data frame.

```
val df2 = Seq((1,"Alice",25,"HR",50000), (2,"Bob",35,"Finance",70000),
(3,"Charlie",28,"HR",60000), (4,"David",45,"IT",90000), (5,"Eva",32,"Finance",80000),
(6,"Frank",29,"IT",75000), (7,"Grace",40,"HR",65000), (8,"Hank",38,"IT",85000),
(9,"Ivy",30,"Finance",72000),
(10,"Jack",50,"HR",78000)).toDF("emp_id","name","age","department","salary")
```

```
cloudera@quickstart:~
File Edit View Search Terminal Help

scala> val df2 = Seq((1,"Alice",25,"HR",50000), (2,"Bob",35,"Finance",70000), (3,"Charlie",28,"HR",60000), (4,"David",45,"IT",90000), (5,"Eva",32,"Finance",80000), (6,"Frank",29,"IT",75000), (7,"Grace",40,"HR",65000), (8,"Hank",38,"IT",85000), (9,"Ivy",30,"Finance",72000), (10,"Jack",50,"HR",78000)).toDF("emp_id","name","age","department","salary")
df2: org.apache.spark.sql.DataFrame = [emp_id: int, name: string, age: int, department: string, salary: int]

scala> df2.show()
+-----+-----+-----+-----+-----+
|emp_id|name|age|department|salary|
+-----+-----+-----+-----+-----+
|1|Alice|25|HR|50000|
|2|Bob|35|Finance|70000|
|3|Charlie|28|HR|60000|
|4|David|45|IT|90000|
|5|Eva|32|Finance|80000|
|6|Frank|29|IT|75000|
|7|Grace|40|HR|65000|
|8|Hank|38|IT|85000|
|9|Ivy|30|Finance|72000|
|10|Jack|50|HR|78000|
+-----+-----+-----+-----+-----+

scala>
```

Step 2: Join the two dataframes using user_id from df and emp_id from df2 and then display it.

```
val joinedDF = df.join( df("user_id") === df2("emp_id"), "left")
joinedDF.show()
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
scala> val joinedDF = df.join(df2, df("user_id") === df2("emp_id"), "left")  
joinedDF: org.apache.spark.sql.DataFrame = [transaction_id: int, user_id: int, amount: int, status: string, emp_id: int, name: string, age: int, department: string, salary: int]  
scala> joinedDF.show()  
+-----+-----+-----+-----+-----+-----+-----+-----+  
|transaction_id|user_id|amount|status|emp_id|name|age|department|salary|  
+-----+-----+-----+-----+-----+-----+-----+-----+  
|101|1|250|Completed|1|Alice|25|HR|50000|  
|102|2|400|Failed|2|Bob|35|Finance|70000|  
|103|3|150|Completed|3|Charlie|28|HR|60000|  
|104|4|600|Completed|4|David|45|IT|90000|  
|105|5|320|Failed|5|Eva|32|Finance|80000|  
|106|1|290|Completed|1|Alice|25|HR|50000|  
|107|2|500|Completed|2|Bob|35|Finance|70000|  
|108|3|120|Failed|3|Charlie|28|HR|60000|  
|109|4|620|Completed|4|David|45|IT|90000|  
|110|5|350|Completed|5|Eva|32|Finance|80000|  
+-----+-----+-----+-----+-----+-----+-----+-----+  
scala>
```

4. Filter rows in the transaction DataFrame where status is Completed and group by user_id.

Step 1: Filter the rows where the status is completed.

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
+-----+-----+-----+-----+-----+-----+-----+-----+  
|108|3|120|Failed|3|Charlie|28|HR|60000|  
|109|4|620|Completed|4|David|45|IT|90000|  
|110|5|350|Completed|5|Eva|32|Finance|80000|  
+-----+-----+-----+-----+-----+-----+-----+-----+  
scala> val completedTransactions = df.filter(df("status") === "Completed")  
completedTransactions: org.apache.spark.sql.DataFrame = [transaction_id: int, user_id: int, amount: int, status: string]  
scala> completedTransactions.show()  
+-----+-----+-----+-----+  
|transaction_id|user_id|amount|status|  
+-----+-----+-----+-----+  
|101|1|250|Completed|  
|103|3|150|Completed|  
|104|4|600|Completed|  
|106|1|290|Completed|  
|107|2|500|Completed|  
|109|4|620|Completed|  
|110|5|350|Completed|  
+-----+-----+-----+-----+  
scala>
```

Step 2: Then from the filtered rows, group the rows based on user_id and show the user_id, count (i.e., number of entries for a particular user) and total amount.

```
scala> val groupedDF = completedTransactions.groupBy("user_id").agg(count("*").alias("completed_count"), sum("amount").alias("total_amount"))  
groupedDF: org.apache.spark.sql.DataFrame = [user_id: int, completed_count: bigint, total_amount: bigint]  
scala> groupedDF.show()  
+-----+-----+-----+  
|user_id|completed_count|total_amount|  
+-----+-----+-----+  
|1|2|540|  
|2|1|500|  
|3|1|150|  
|4|2|1220|  
|5|1|350|  
+-----+-----+-----+  
scala>
```

5. Cache the employee DataFrame and use it in multiple queries to improve performance.

Step 1: Cache the employee data frame that was created before.

```
df2.cache()
```

Step 2: Perform multiple queries with that data frame. Here I have:

Displayed the data: df2.show()

Filtered the data to display only the rows where department is HR: val hrEmployees = df2.filter(df2("department") === "HR")

```
scala> df2.cache()
res9: df2.type = [emp_id: int, name: string, age: int, department: string, salary: int]

scala> df2.show()
+-----+-----+-----+-----+-----+
|emp_id|  name|age|department|salary|
+-----+-----+-----+-----+
|    1| Alice| 25|        HR| 50000|
|    2|  Bob| 35|    Finance| 70000|
|    3|Charlie| 28|        HR| 60000|
|    4| David| 45|        IT| 90000|
|    5|  Eva| 32|    Finance| 80000|
|    6| Frank| 29|        IT| 75000|
|    7| Grace| 40|        HR| 65000|
|    8|  Hank| 38|        IT| 85000|
|    9|  Ivy| 30|    Finance| 72000|
|   10| Jack| 50|        HR| 78000|
+-----+-----+-----+-----+

scala> val hrEmployees = df2.filter(df2("department") === "HR")
hrEmployees: org.apache.spark.sql.DataFrame = [emp_id: int, name: string, age: int, department: string, salary: int]

scala> hrEmployees.show()
+-----+-----+-----+-----+-----+
|emp_id|  name|age|department|salary|
+-----+-----+-----+-----+
|    1| Alice| 25|        HR| 50000|
|    3|Charlie| 28|        HR| 60000|
|    7| Grace| 40|        HR| 65000|
|   10| Jack| 50|        HR| 78000|
+-----+-----+-----+-----+

cloudera@quickstart:~
```

Computed the average salary of the people working in the Finance Department:

```
val avgFinanceSalary = df2.filter(df2("department") === "Finance").agg(avg("salary").alias("average_salary"))
```

```
avgFinanceSalary.show()
```

```
scala> val avgFinanceSalary = df2.filter(df2("department") === "Finance").agg(avg("salary").alias("average_salary"))
avgFinanceSalary: org.apache.spark.sql.DataFrame = [average_salary: double]

scala> avgFinanceSalary.show()
+-----+
|average_salary|
+-----+
|      74000.0|
+-----+

scala>
```